


B as a single-ordering benchmark: native initialization

MPCurve 0.3.4 | princurve 2.1.6 | GPy 1.13.2; default initialization, extended optimization budgets.



All methods fit one ordering to the same 300 samples and 12 features; each uses its native PCA start. Blue: initial ranks; orange: final ranks; gray lines link the same sample. Global reversal is aligned away.

MPCurve: up to 2000 sweeps, tol = 1e-6. P-curve: maxit = 1000, thresh = 1e-6; default df = 5.

Both GP models use native optimization with up to three 2000-evaluation blocks; Bayesian GPLVM uses the default 10 inducing points.

The two noiseless GP fits reached their evaluation budgets; those endpoints are not reported as converged.